

## ▼ Random forest using grid search

## ▼ Methodology

- **Data Cleaning:** Checking for null values and based on their number either dropping them or replacing with mean, median, mode based on the type and description of data. Dropping discrete and categorical variables that have highly skewed histograms.
- **Data Visualization:** This step helps understand the data in a visually. We can understand normality of the data as well. This helps us to decide whether to normalize the data. In case of categorical variables it also helps in feature selection.
- **Feature Selection:** Based on the Pearson correlation between the labeled column and rest of the features. In general, a very great correlation should have an absolute value greater than 0.75. When the labeled column is depended on multiple columns, the correlation with one column may be less. But combined features may have higher effect.
- **Train Test Split:** We split the data into 80:20 ratio for training testing respectively.
- **Model Selection:** Based on the data visualization and data correlation, we need to select a model that would best suit. Here we need to use XGBOOST.
- **Evaluation:** In this case we are using RMSE, R2 Score to determine the accuracy of the predicting model.

## ▼ importing libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

## ▼ Reading data

```
df=pd.read_csv(r"drive/My Drive/Asteroid_Updated.csv")
```

```
↳ /usr/local/lib/python3.6/dist-packages/IPython/core/interactiveshell.py:2718: DtypeWarning: Columns (0,10,15,16,23,24) have mixed types.Specify dtype option on import or set low_memory=False
interactivity=interactivity, compiler=compiler, result=result)
```

## ▼ Null value percentages

```
Null=[]
for i in df:
    Null.append((i,df[i].isna().mean()*100))
Null=pd.DataFrame(Null,columns=['class', 'per'])
Null
```

```
↳
```

	class	per
0	name	97.383990
1	a	0.000238
2	e	0.000000
3	i	0.000000
4	om	0.000000
5	w	0.000000
6	q	0.000000
7	ad	0.000715
8	per_y	0.000119
9	data_arc	1.842770
10	condition_code	0.103249
11	n_obs_used	0.000000
12	H	0.320228
13	neo	0.000715
14	pha	1.958048
15	diameter	83.609181
16	extent	99.997856
17	albedo	83.755302
18	rot_per	97.761619
19	GM	99.998333
20	BV	99.878411
21	UB	99.883413
22	IR	99.999881
23	spec_B	99.801599
24	spec_T	99.883294
25	G	99.985829
26	moid	1.958048

▼ Columns with percentage of null values Greater than 85%

```
l=Null[Null['per']>85]['class']
Null[Null['per']>85]
```

🔗

	class	per
0	name	97.383990
16	extent	99.997856
18	rot_per	97.761619
19	GM	99.998333
20	BV	99.878411


These columns will have very less impact on the decision since most of their values are null.  
We can drop them

```
df=df.drop(1,axis=1)
df
```

	a	e	i	om	w	q	ad	per_y	data_arc	condition_code	n_obs_used	H	neo	pha	diameter	albedo	moid	class	n
0	2.769165	0.076009	10.594067	80.305532	73.597694	2.558684	2.979647	4.608202	8822.0	0	1002	3.340	N	N	939.4	0.0900	1.594780	MBA	0.213885
1	2.772466	0.230337	34.836234	173.080063	310.048857	2.133865	3.411067	4.616444	72318.0	0	8490	4.130	N	N	545	0.1010	1.233240	MBA	0.213503
2	2.669150	0.256942	12.988919	169.852760	248.138626	1.983332	3.354967	4.360814	72684.0	0	7104	5.330	N	N	246.596	0.2140	1.034540	MBA	0.226019
3	2.361418	0.088721	7.141771	103.810804	150.728541	2.151909	2.570926	3.628837	24288.0	0	9325	3.200	N	N	525.4	0.4228	1.139480	MBA	0.271609
4	2.574249	0.191095	5.366988	141.576605	358.687607	2.082324	3.066174	4.130323	63507.0	0	2916	6.850	N	N	106.699	0.2740	1.095890	MBA	0.238632
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
839709	2.812945	0.664688	4.695700	183.310012	234.618352	0.943214	4.682676	4.717914	17298.0	0	118	20.400	Y	Y	NaN	NaN	0.032397	APO	0.208911
839710	2.645238	0.259376	12.574937	1.620020	339.568072	1.959126	3.331350	4.302346	16.0	9	15	17.507	N	N	NaN	NaN	0.956145	MBA	0.229090
839711	2.373137	0.202053	0.732484	176.499082	198.026527	1.893638	2.852636	3.655884	5.0	9	6	18.071	N	N	NaN	NaN	0.893896	MBA	0.269600
839712	2.260404	0.258348	9.661947	204.512448	148.496988	1.676433	2.844376	3.398501	10.0	9	13	18.060	N	N	NaN	NaN	0.680220	MBA	0.290018
839713	2.546442	0.287672	5.356238	70.709555	273.483265	1.813901	3.278983	4.063580	11.0	9	11	17.406	N	N	NaN	NaN	0.815280	MBA	0.242551

839714 rows × 21 columns

```
df.dtypes
```

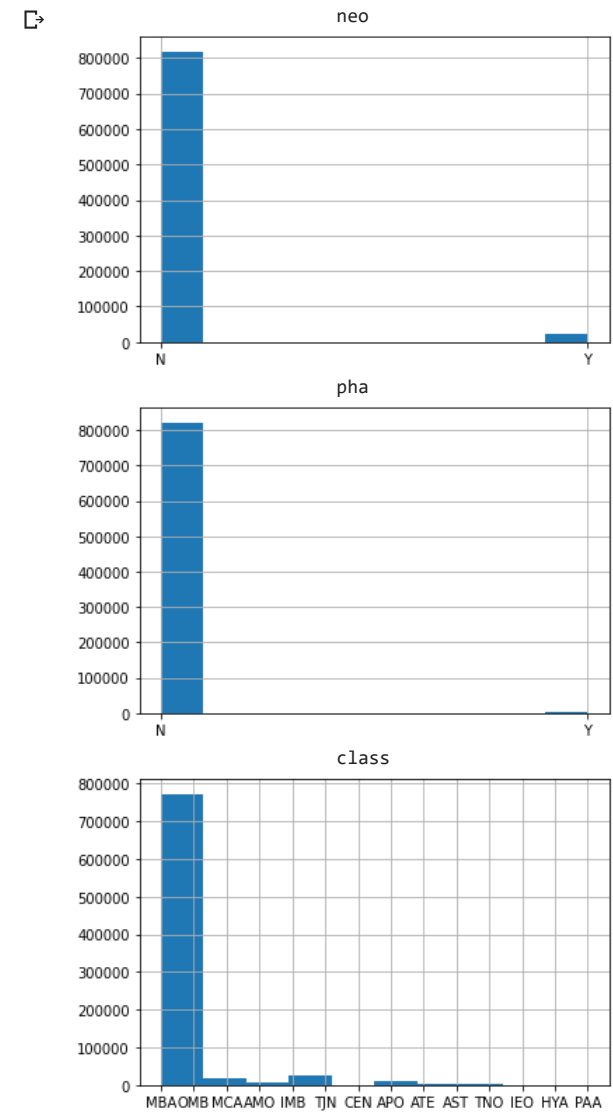


a	float64
e	float64
i	float64

▼ Catagorical columns with skewed *distribution*

q	float64
---	---------

```
k=['neo','pha','class']
for i in k:
    print("\t\t\t",i)
    df[i].hist()
    plt.show()
```



Since the most values have only one catagory so, the impact of rest of the catagories is not significant, and dropping them will not effect the decision variable.

```
df=df.drop(k,axis=1)
```

df.dtypes

```
a          float64
e          float64
i          float64
om         float64
w          float64
q          float64
ad         float64
per_y      float64
data_arc   float64
condition_code  object
n_obs_used  int64
H          float64
diameter   object
albedo     float64
moid       float64
n          float64
per        float64
ma         float64
dtype: object
```

#### ▼ Handling diameter columns data

```
df.diameter.unique()
```

```
array(['939.4', '545', '246.596', ..., 0.122, 0.6509999999999999, 1.077],
      dtype=object)
```

The data is in both numeric and catagorical values, we need to convert it to int data type

```
df['diameter']= df.diameter.astype(float,errors='ignore')
```

```
df.diameter.unique()
```

```
array([9.394000e+02, 5.45000e+02, 2.46596e+02, ..., 1.22000e-01,
        6.51000e-01, 1.07700e+00])
```

#### ▼ Droping null values

```
df.isna().sum()/len(df)*100
```



```

a                0.000238
~ ~ ~ ~ ~ ~ ~ ~
```

there are null values in predicting column and it is a high number filling them increases the error in the data so, dropping them is only the option.

```

w                a 0000000
df = df[df['diameter'].notna()]
```

```

nan v           a 0000110
fill the remaining with median value
```

```

n obs used      0.000000
df=df.fillna(df.mean())
df.isna().sum()
```

```

[> a                0
   e                0
   i                0
   om               0
   w                0
   q                0
   ad               0
   per_y            0
   data_arc         0
   condition_code   0
   n_obs_used       0
   H                0
   diameter         0
   albedo           0
   moid             0
   n                0
   per              0
   ma               0
   dtype: int64
```

▼ Handling condition\_code columns data

```
df['condition_code'].unique()
```

```

[> array([0, 1, 3, 2, '0', '1', '2', '3', '4', '5', '9', '7', 5.0, 6.0, 4.0,
        7.0, 9.0, 8.0, '8', '6'], dtype=object)
```

```
df['condition_code']= df.condition_code.astype(int)
```

```
df['condition_code'].unique()
```

```

[> array([0, 1, 3, 2, 4, 5, 9, 7, 6, 8])
```


▼ Feature

```
import seaborn as sns
plt.figure(figsize=(30,20))
sns.heatmap(df.corr(),annot = True,cmap="jet")
```

```

[>
```

```
import pandas.util.testing as tm
<matplotlib.axes._subplots.AxesSubplot at 0x7f3d67f18ef0>
```



abs(df.corr()['diameter']\*100).sort\_values(ascending=False)





```
diameter      100.000000
H              56.849287
data_arc       49.157968
n_obs_used     20.574724
```

15 highly correlated features

```
n      20.102275
```

```
l=dict(abs(df.corr()[ 'diameter']*100).sort_values(ascending=False)[:3]).keys()
l
```

```
dict_keys(['diameter', 'H', 'data_arc', 'n_obs_used', 'moid', 'q', 'n', 'a', 'albedo', 'ad', 'condition_code', 'i', 'e', 'per', 'per_y'])
e      4.913337
```

```
df=df[l]
```

```
ma      0.965894
```

## Train Test Split

```
name, diameter, dtype, float64
```

```
cols = [col for col in df.columns if col not in ["diameter"]]
X = df[cols]
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df, df['diameter'], test_size=0.2)
```

## XGBOOST Model

```
from xgboost import XGBRegressor
```

```
model = XGBRegressor(max_depth=8,n_jobs=6,booster='dart')
model.fit(X_train, y_train)
```

```
[19:25:16] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.
XGBRegressor(base_score=0.5, booster='dart', colsample_bylevel=1,
             colsample_bynode=1, colsample_bytree=1, gamma=0,
             importance_type='gain', learning_rate=0.1, max_delta_step=0,
             max_depth=8, min_child_weight=1, missing=None, n_estimators=100,
             n_jobs=6, nthread=None, objective='reg:linear', random_state=0,
             reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
             silent=None, subsample=1, verbosity=1)
```

## R2 Score on Test and Train set

```
model.score(X_train, y_train)
```

```
0.9999974603797344
```

```
model.score(X_test, y_test)
```

```
0.9205378495840809
```