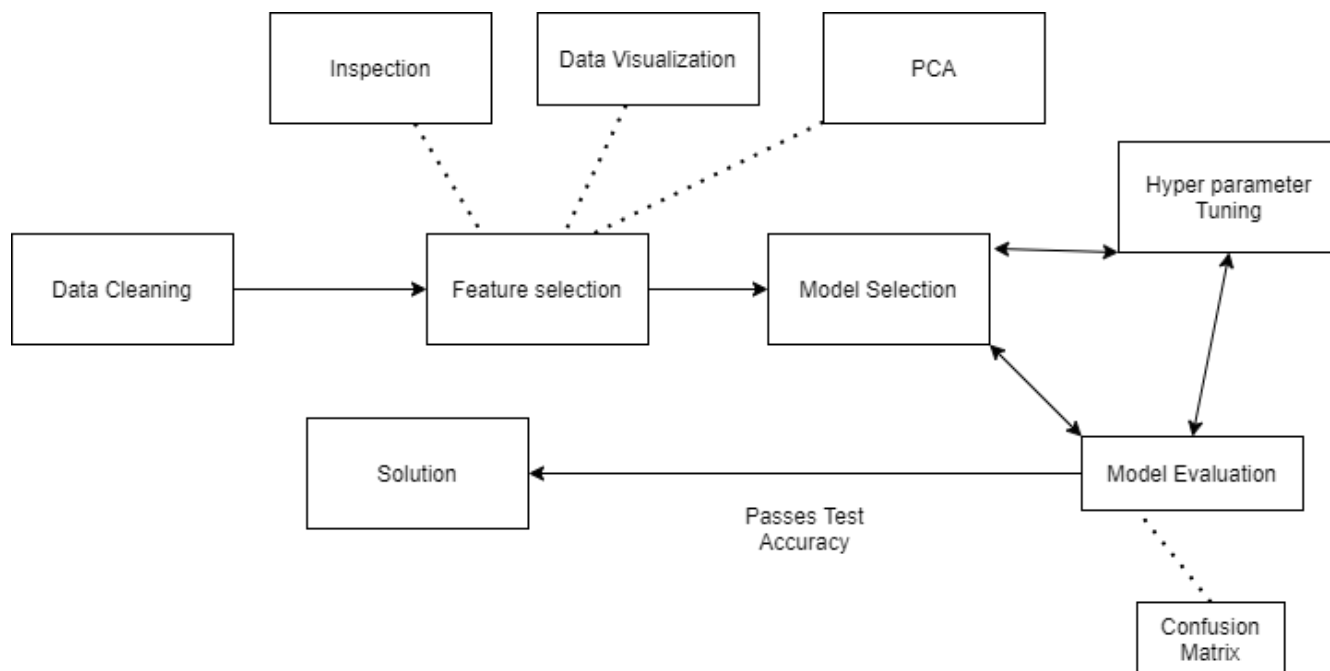


# SVM

## Methodology

- ##### **Data Cleaning:** Checking for null values and based on their number either dropping them or replacing with mean, median, mode based on the type and description of data. Dropping discrete and categorical variables that have highly skewed histograms.
- ##### **Data Visualization:** This step helps understand the data in a visually. We can understand normality of the data as well. This helps us to decide whether to normalize the data. In case of categorical variables it also helps in feature selection.
- ##### **Feature Selection:** Based on the Pearson correlation between the labeled column and rest of the features. In general, a very great correlation should have an absolute value greater than 0.75. When the labeled column is depended on multiple columns, the correlation with one column may be less. But combined features may have higher effect.
- ##### **Train Test Split:** We split the data into 80:20 ratio for training testing respectively.
- ##### **Model Selection:** Based on the data visualization and data correlation, we need to select a model that would best suit. Here we need to use SVM.
- ##### **Evaluation:** In this case we are using F1 Score to determine the accuracy of the predicting model.



In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
```

In [2]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

In [3]:

```
df = pd.read_excel("drive/My Drive/voice.xlsx")
```

In [4]:

```
df
```

Out[4]:

	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt	sp.ent	sfm	mode	cen
0	0.059781	0.064241	0.032027	0.015071	0.090193	0.075122	12.863462	274.402906	0.893369	0.491918	0.000000	0.05
1	0.066009	0.067310	0.040229	0.019414	0.092666	0.073252	22.423285	634.613855	0.892193	0.513724	0.000000	0.06
2	0.077316	0.083829	0.036718	0.008701	0.131908	0.123207	30.757155	1024.927705	0.846389	0.478905	0.000000	0.07
3	0.151228	0.072111	0.158011	0.096582	0.207955	0.111374	1.232831	4.177296	0.963322	0.727232	0.083878	0.15
4	0.135120	0.079146	0.124656	0.078720	0.206045	0.127325	1.101174	4.333713	0.971955	0.783568	0.104261	0.13
...	...	...	...	...	...	...	...	...	...	...	...	...
3163	0.131884	0.084734	0.153707	0.049285	0.201144	0.151859	1.762129	6.630383	0.962934	0.763182	0.200836	0.13
3164	0.116221	0.089221	0.076758	0.042718	0.204911	0.162193	0.693730	2.503954	0.960716	0.709570	0.013683	0.11
3165	0.142056	0.095798	0.183731	0.033424	0.224360	0.190936	1.876502	6.604509	0.946854	0.654196	0.008006	0.14
3166	0.143659	0.090628	0.184976	0.043508	0.219943	0.176435	1.591065	5.388298	0.950436	0.675470	0.212202	0.14
3167	0.165509	0.092884	0.183044	0.070072	0.250827	0.180756	1.705029	5.769115	0.938829	0.601529	0.267702	0.16

3168 rows x 21 columns



In [5]:

```
df.isna().sum()
```

Out[5]:

```
meanfreq    0
sd           0
median       0
Q25          0
Q75          0
IQR          0
skew         0
kurt         0
sp.ent       0
sfm          0
mode         0
centroid     0
meanfun      0
minfun       0
maxfun       0
meandom      0
mindom       0
maxdom       0
dfrange      0
modindx      0
label        0
dtype: int64
```

NO Null values

In [7]:

```
df.dtypes
```

Out[7]:

```
meanfreq    float64
sd           float64
median       float64
Q25          float64
Q75          float64
IQR          float64
skew         float64
kurt         float64
sp.ent       float64
sfm          float64
```

```
mode          float64
centroid      float64
meanfun       float64
minfun        float64
maxfun        float64
meandom       float64
mindom        float64
maxdom        float64
dfrange       float64
modindx       float64
label         object
dtype: object
```

In [8]:

```
df["label"]=df["label"].replace('male',1)
df["label"]=df["label"].replace('female',2)
df["label"].unique()
```

Out[8]:

```
array([1, 2])
```

In [9]:

```
df.dtypes
```


Out[9]:

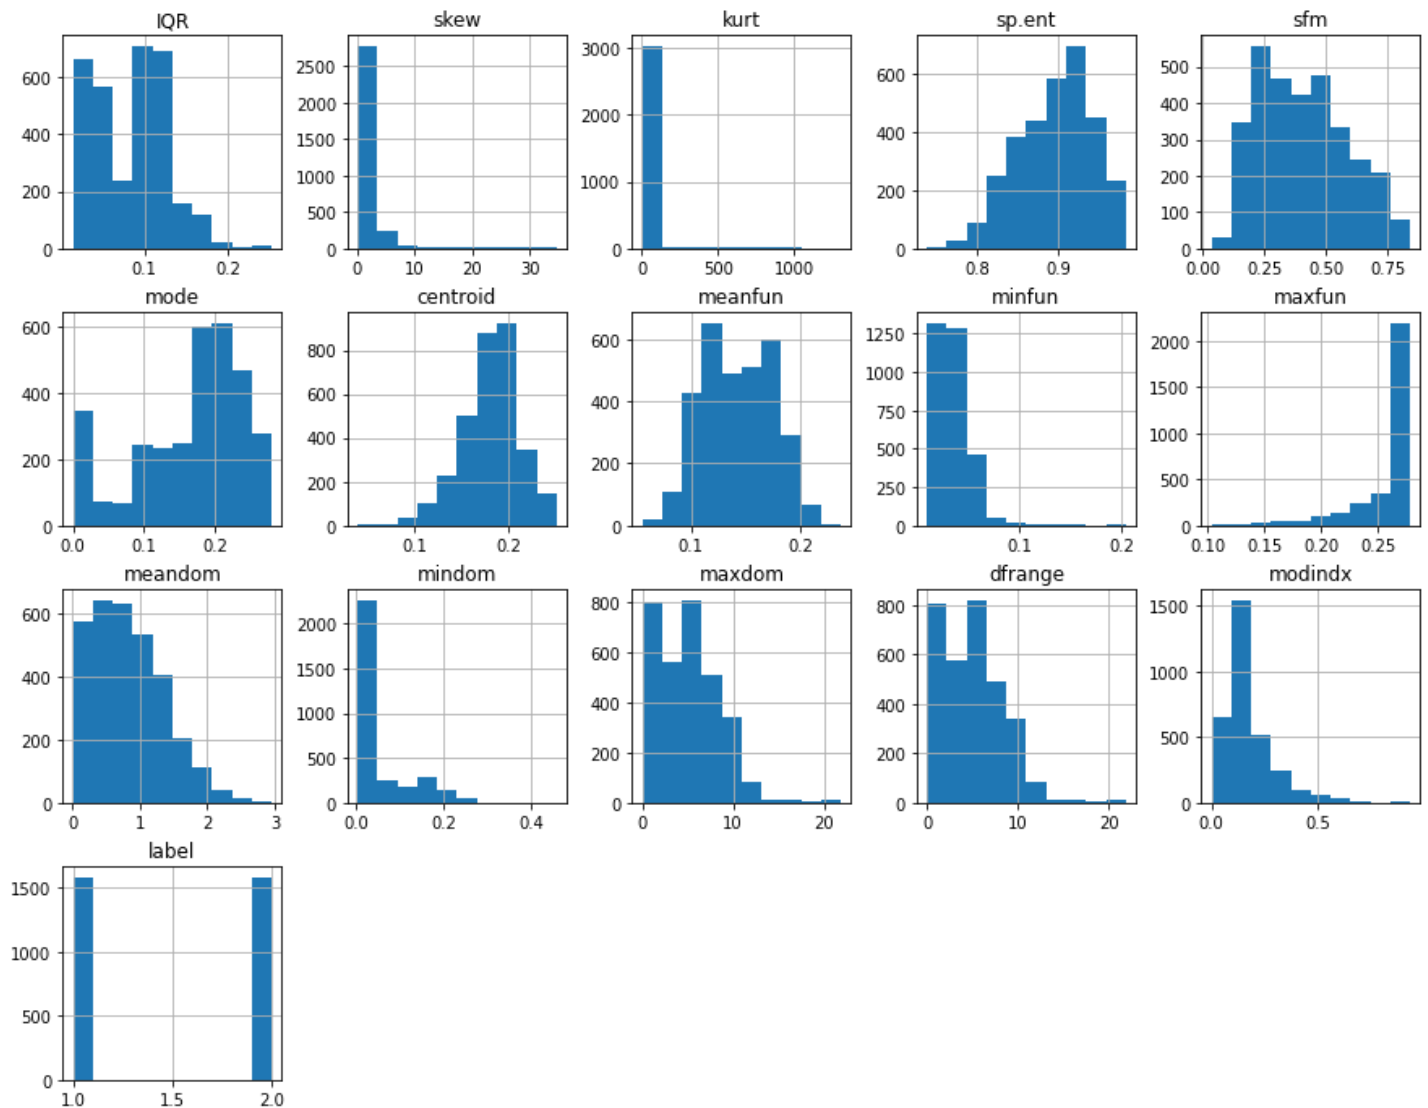
```
meanfreq      float64
sd            float64
median        float64
Q25          float64
Q75          float64
IQR          float64
skew          float64
kurt          float64
sp.ent        float64
sfm           float64
mode          float64
centroid      float64
meanfun       float64
minfun        float64
maxfun        float64
meandom       float64
mindom        float64
maxdom        float64
dfrange       float64
modindx       float64
label         int64
dtype: object
```

In [12]:

```
c=list(df)
fig=plt.figure(figsize=(15,15))
ax=fig.gca()
df[c].hist(ax=ax)
plt.show()
```

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:4: UserWarning: To output multiple subplots, the figure containing the passed axes is being cleared after removing the cwd from sys.path.





Kurt and skew has most of the values as only one value so removing it will not make any difference in decision making.

The label data is balanced, so no need for sampling.

In [13]:

```
df=df.drop(["kurt","skew"],axis=1)
df
```

Out[13]:

	meanfreq	sd	median	Q25	Q75	IQR	sp.ent	sfm	mode	centroid	meanfun	minfun
0	0.059781	0.064241	0.032027	0.015071	0.090193	0.075122	0.893369	0.491918	0.000000	0.059781	0.084279	0.015702
1	0.066009	0.067310	0.040229	0.019414	0.092666	0.073252	0.892193	0.513724	0.000000	0.066009	0.107937	0.015826
2	0.077316	0.083829	0.036718	0.008701	0.131908	0.123207	0.846389	0.478905	0.000000	0.077316	0.098706	0.015656
3	0.151228	0.072111	0.158011	0.096582	0.207955	0.111374	0.963322	0.727232	0.083878	0.151228	0.088965	0.017798
4	0.135120	0.079146	0.124656	0.078720	0.206045	0.127325	0.971955	0.783568	0.104261	0.135120	0.106398	0.016931
...	...	...	...	...	...	...	...	...	...	...	...	...
3163	0.131884	0.084734	0.153707	0.049285	0.201144	0.151859	0.962934	0.763182	0.200836	0.131884	0.182790	0.083770
3164	0.116221	0.089221	0.076758	0.042718	0.204911	0.162193	0.960716	0.709570	0.013683	0.116221	0.188980	0.034409
3165	0.142056	0.095798	0.183731	0.033424	0.224360	0.190936	0.946854	0.654196	0.008006	0.142056	0.209918	0.039506
3166	0.143659	0.090628	0.184976	0.043508	0.219943	0.176435	0.950436	0.675470	0.212202	0.143659	0.172375	0.034483
3167	0.165509	0.092884	0.183044	0.070072	0.250827	0.180756	0.938829	0.601529	0.267702	0.165509	0.185607	0.062257


3168 rows x 19 columns

In [20]:

```
import seaborn as sns
ax=plt.figure(figsize=(15,15))
sns.boxplot(data=df)
```

Out[20]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f36455f20b8>




In [21]:

```
df_norm = (df - df.min()) / (df.max() - df.min())
ax=plt.figure(figsize=(15,15))
sns.boxplot(data=df_norm) # boxplot for normalized data
```

Out[21]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f36457f9898>





In [22]:

```
cols = [col for col in df.columns if col not in ["label"]]
X = df[cols]
y = df["label"]
```

In [23]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

## Logistic

In [24]:

```
from sklearn.linear_model import LogisticRegression
clf = LogisticRegression(random_state=0).fit(X_train, y_train)
```

/usr/local/lib/python3.6/dist-packages/sklearn/linear\_model/\_logistic.py:940: Convergence Warning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

In [25]:

```
from sklearn.metrics import classification_report
p_train = clf.predict(X_train)
p_test = clf.predict(X_test)
```

In [27]:

```
print(classification_report(y_train, p_train, target_names=["male", "female"])) # report on train data
```

	precision	recall	f1-score	support
male	0.87	0.98	0.92	1279
female	0.97	0.85	0.90	1255
accuracy			0.91	2534
macro avg	0.92	0.91	0.91	2534
weighted avg	0.92	0.91	0.91	2534

In [28]:

```
print(classification_report(y_test, p_test, target_names=["male", "female"])) # report on test data
```

	precision	recall	f1-score	support
male	0.83	0.98	0.90	305
female	0.97	0.82	0.89	329
accuracy			0.90	634
macro avg	0.90	0.90	0.90	634
weighted avg	0.91	0.90	0.90	634

In [31]:

```
from sklearn.metrics import confusion_matrix
confm=confusion_matrix(y_test, p_test, labels=[1,2]) # confusion matrix on test data
confm
```

Out[31]:

```
array([[298,  7],
       [ 59, 270]])
```

**There is less variance in the model**

## SVM

In [33]:

```
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
param_grid = {'C': [0.01, 0.1, 1, 10], 'kernel': ('rbf', 'linear')}
classifier = SVC()
grid_search = GridSearchCV(estimator=classifier, param_grid=param_grid, scoring='accuracy', n_jobs=-1, verbose=42)
grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 8 candidates, totalling 40 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 1 tasks | elapsed: 1.7s
[Parallel(n_jobs=-1)]: Done 2 tasks | elapsed: 1.7s
[Parallel(n_jobs=-1)]: Done 3 tasks | elapsed: 2.2s
[Parallel(n_jobs=-1)]: Done 4 tasks | elapsed: 2.2s
[Parallel(n_jobs=-1)]: Done 5 tasks | elapsed: 2.4s
[Parallel(n_jobs=-1)]: Done 6 tasks | elapsed: 2.6s
[Parallel(n_jobs=-1)]: Done 7 tasks | elapsed: 2.6s
[Parallel(n_jobs=-1)]: Done 8 tasks | elapsed: 2.8s
```

```
[Parallel(n_jobs=-1)]: Done    9 tasks      | elapsed:    2.9s
[Parallel(n_jobs=-1)]: Done   10 tasks      | elapsed:    3.0s
[Parallel(n_jobs=-1)]: Done   11 tasks      | elapsed:    3.3s
[Parallel(n_jobs=-1)]: Done   12 tasks      | elapsed:    3.4s
[Parallel(n_jobs=-1)]: Done   13 tasks      | elapsed:    3.6s
[Parallel(n_jobs=-1)]: Done   14 tasks      | elapsed:    3.8s
[Parallel(n_jobs=-1)]: Done   15 tasks      | elapsed:    4.0s
[Parallel(n_jobs=-1)]: Done   16 tasks      | elapsed:    4.1s
[Parallel(n_jobs=-1)]: Done   17 tasks      | elapsed:    4.3s
[Parallel(n_jobs=-1)]: Done   18 tasks      | elapsed:    4.4s
[Parallel(n_jobs=-1)]: Done   19 tasks      | elapsed:    4.6s
[Parallel(n_jobs=-1)]: Done   20 tasks      | elapsed:    4.6s
[Parallel(n_jobs=-1)]: Done   21 tasks      | elapsed:    4.9s
[Parallel(n_jobs=-1)]: Done   22 tasks      | elapsed:    5.0s
[Parallel(n_jobs=-1)]: Done   23 tasks      | elapsed:    5.3s
[Parallel(n_jobs=-1)]: Done   24 tasks      | elapsed:    5.4s
[Parallel(n_jobs=-1)]: Done   25 tasks      | elapsed:    5.7s
[Parallel(n_jobs=-1)]: Done   26 tasks      | elapsed:    5.8s
[Parallel(n_jobs=-1)]: Done   27 tasks      | elapsed:    6.1s
[Parallel(n_jobs=-1)]: Done   28 tasks      | elapsed:    6.3s
[Parallel(n_jobs=-1)]: Done   29 tasks      | elapsed:    6.5s
[Parallel(n_jobs=-1)]: Done   30 tasks      | elapsed:    6.8s
[Parallel(n_jobs=-1)]: Done   31 tasks      | elapsed:    6.9s
[Parallel(n_jobs=-1)]: Done   32 tasks      | elapsed:    7.1s
[Parallel(n_jobs=-1)]: Done   33 tasks      | elapsed:    7.2s
[Parallel(n_jobs=-1)]: Done   34 tasks      | elapsed:    7.5s
[Parallel(n_jobs=-1)]: Done   35 tasks      | elapsed:    7.6s
[Parallel(n_jobs=-1)]: Done   36 tasks      | elapsed:    8.2s
[Parallel(n_jobs=-1)]: Done   37 tasks      | elapsed:    8.2s
[Parallel(n_jobs=-1)]: Done  38 out of  40 | elapsed:    8.9s remaining:    0.5s
[Parallel(n_jobs=-1)]: Done  40 out of  40 | elapsed:    9.3s remaining:    0.0s
[Parallel(n_jobs=-1)]: Done  40 out of  40 | elapsed:    9.3s finished
```

Out[33]:

```
GridSearchCV(cv=None, error_score=nan,
             estimator=SVC(C=1.0, break_ties=False, cache_size=200,
                           class_weight=None, coef0=0.0,
                           decision_function_shape='ovr', degree=3,
                           gamma='scale', kernel='rbf', max_iter=-1,
                           probability=False, random_state=None, shrinking=True,
                           tol=0.001, verbose=False),
             iid='deprecated', n_jobs=-1,
             param_grid={'C': [0.01, 0.1, 1, 10], 'kernel': ('rbf', 'linear')},
             pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
             scoring='accuracy', verbose=42)
```

In [34]:

```
p_train = grid_search.predict(X_train)
p_test = grid_search.predict(X_test)
```

In [35]:

```
print(classification_report(y_train, p_train, target_names=["male", "female"])) # report
on train data
```

	precision	recall	f1-score	support
male	0.96	0.98	0.97	1279
female	0.98	0.96	0.97	1255
accuracy			0.97	2534
macro avg	0.97	0.97	0.97	2534
weighted avg	0.97	0.97	0.97	2534

In [36]:

```
print(classification_report(y_test, p_test, target_names=["male", "female"])) # report on
test data
```



	precision	recall	f1-score	support
male	0.96	0.96	0.96	305
female	0.97	0.97	0.97	329
accuracy			0.97	634
macro avg	0.97	0.97	0.97	634
weighted avg	0.97	0.97	0.97	634

In [37]:

```
from sklearn.metrics import confusion_matrix
confm=confusion_matrix(y_test, p_test, labels=[1,2]) # confusion matrix on test data
confm
```

Out[37]:

```
array([[294, 11],
       [ 11, 318]])
```

**The error has reduced and the accuracy increased in both classes. The model has better bias and variance scores than Logistic regression.**

In [ ]: